

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df = pd.read_csv("50_Startups.csv")
df.head()
```

```
Out[2]:
```

| | R&D Spend | Administration | Marketing Spend | State | Profit |
|---|-----------|----------------|-----------------|------------|-----------|
| 0 | 165349.20 | 136897.80 | 471784.10 | New York | 192261.83 |
| 1 | 162597.70 | 151377.59 | 443898.53 | California | 191792.06 |
| 2 | 153441.51 | 101145.55 | 407934.54 | Florida | 191050.39 |
| 3 | 144372.41 | 118671.85 | 383199.62 | New York | 182901.99 |
| 4 | 142107.34 | 91391.77 | 366168.42 | Florida | 166187.94 |

```
In [3]: df.info()

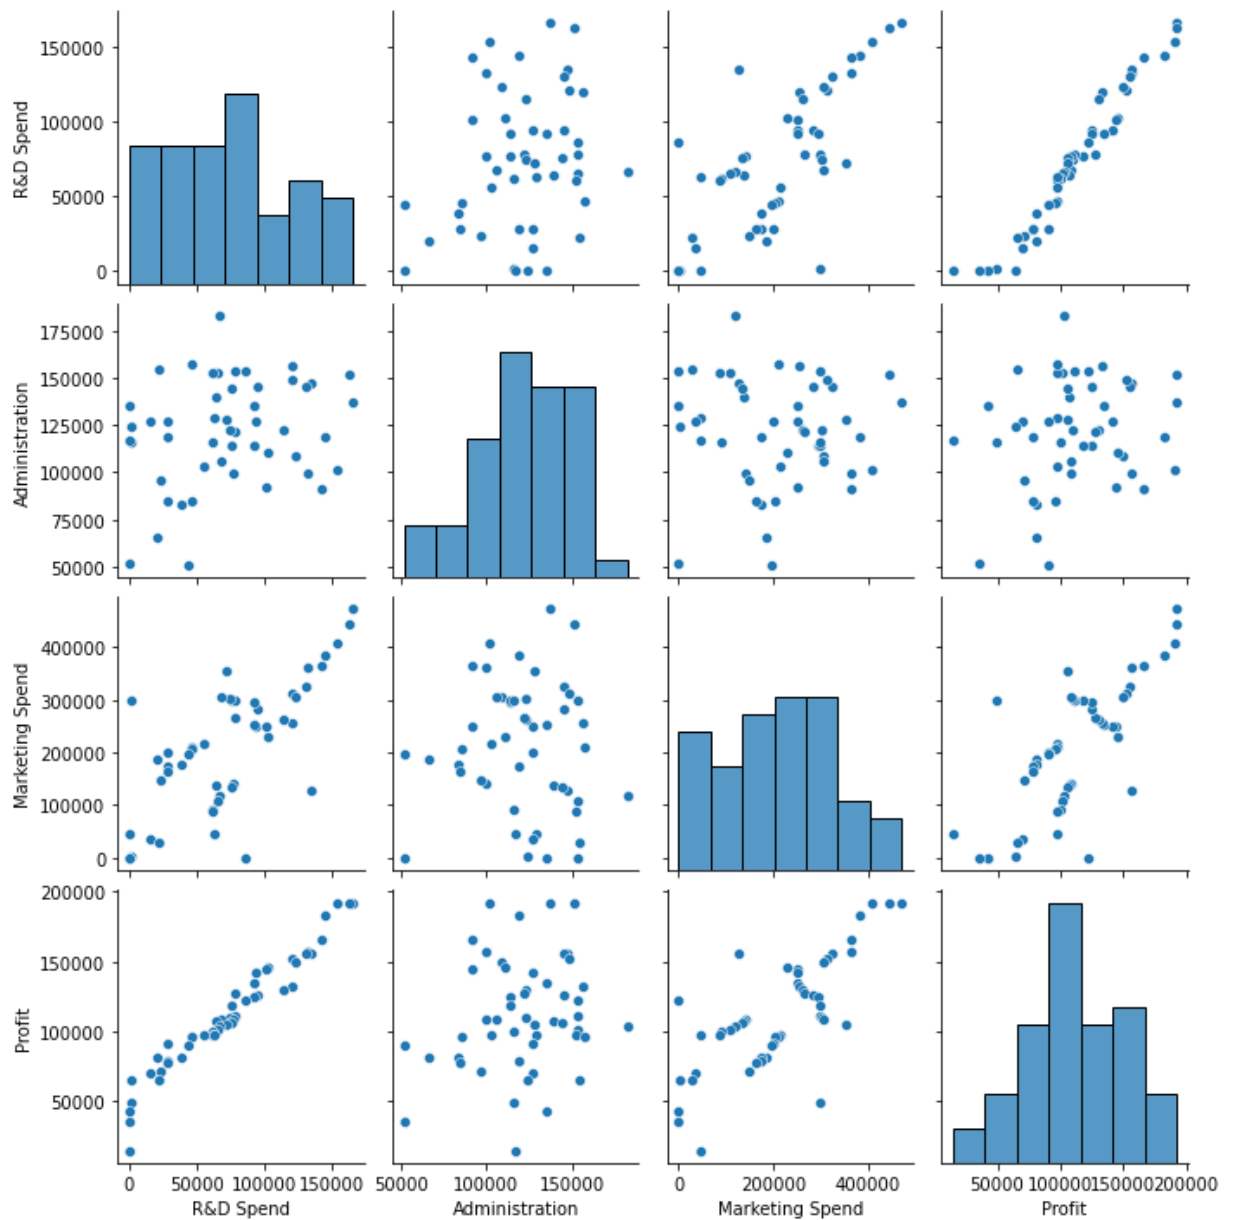
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   R&D Spend              50 non-null     float64
 1   Administration         50 non-null     float64
 2   Marketing Spend        50 non-null     float64
 3   State                  50 non-null     object  
 4   Profit                 50 non-null     float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

```
In [4]: df.isna().sum()
```

```
Out[4]: R&D Spend          0
Administration          0
Marketing Spend          0
State                   0
Profit                  0
dtype: int64
```

In [5]: `sns.pairplot(df)`

Out[5]: `<seaborn.axisgrid.PairGrid at 0x6ab0e37f0>`



In [6]: `x= df.iloc[:, :-1].values # [row:row, col:col]`
`y =df.iloc[:, -1].values`

In [7]: `from sklearn.compose import ColumnTransformer`
`from sklearn.preprocessing import OneHotEncoder`

In [8]: `ct = ColumnTransformer(transformers=[("encoder", OneHotEncoder(), [3])], remainder='passthrough')`
`x = np.array(ct.fit_transform(x))`

In [9]: x

```
Out[9]: array([[0.0, 0.0, 1.0, 165349.2, 136897.8, 471784.1],
 [1.0, 0.0, 0.0, 162597.7, 151377.59, 443898.53],
 [0.0, 1.0, 0.0, 153441.51, 101145.55, 407934.54],
 [0.0, 0.0, 1.0, 144372.41, 118671.85, 383199.62],
 [0.0, 1.0, 0.0, 142107.34, 91391.77, 366168.42],
 [0.0, 0.0, 1.0, 131876.9, 99814.71, 362861.36],
 [1.0, 0.0, 0.0, 134615.46, 147198.87, 127716.82],
 [0.0, 1.0, 0.0, 130298.13, 145530.06, 323876.68],
 [0.0, 0.0, 1.0, 120542.52, 148718.95, 311613.29],
 [1.0, 0.0, 0.0, 123334.88, 108679.17, 304981.62],
 [0.0, 1.0, 0.0, 101913.08, 110594.11, 229160.95],
 [1.0, 0.0, 0.0, 100671.96, 91790.61, 249744.55],
 [0.0, 1.0, 0.0, 93863.75, 127320.38, 249839.44],
 [1.0, 0.0, 0.0, 91992.39, 135495.07, 252664.93],
 [0.0, 1.0, 0.0, 119943.24, 156547.42, 256512.92],
 [0.0, 0.0, 1.0, 114523.61, 122616.84, 261776.23],
 [1.0, 0.0, 0.0, 78013.11, 121597.55, 264346.06],
 [0.0, 0.0, 1.0, 94657.16, 145077.58, 282574.31],
 [0.0, 1.0, 0.0, 91749.16, 114175.79, 294919.57],
 [0.0, 0.0, 1.0, 86419.7, 153514.11, 0.0],
 [1.0, 0.0, 0.0, 76253.86, 113867.3, 298664.47],
 [0.0, 0.0, 1.0, 78389.47, 153773.43, 299737.29],
 [0.0, 1.0, 0.0, 73994.56, 122782.75, 303319.26],
 [0.0, 1.0, 0.0, 67532.53, 105751.03, 304768.73],
 [0.0, 0.0, 1.0, 77044.01, 99281.34, 140574.81],
 [1.0, 0.0, 0.0, 64664.71, 139553.16, 137962.62],
 [0.0, 1.0, 0.0, 75328.87, 144135.98, 134050.07],
 [0.0, 0.0, 1.0, 72107.6, 127864.55, 353183.81],
 [0.0, 1.0, 0.0, 66051.52, 182645.56, 118148.2],
 [0.0, 0.0, 1.0, 65605.48, 153032.06, 107138.38],
 [0.0, 1.0, 0.0, 61994.48, 115641.28, 91131.24],
 [0.0, 0.0, 1.0, 61136.38, 152701.92, 88218.23],
 [1.0, 0.0, 0.0, 63408.86, 129219.61, 46085.25],
 [0.0, 1.0, 0.0, 55493.95, 103057.49, 214634.81],
 [1.0, 0.0, 0.0, 46426.07, 157693.92, 210797.67],
 [0.0, 0.0, 1.0, 46014.02, 85047.44, 205517.64],
 [0.0, 1.0, 0.0, 28663.76, 127056.21, 201126.82],
 [1.0, 0.0, 0.0, 44069.95, 51283.14, 197029.42],
 [0.0, 0.0, 1.0, 20229.59, 65947.93, 185265.1],
 [1.0, 0.0, 0.0, 38558.51, 82982.09, 174999.3],
 [1.0, 0.0, 0.0, 28754.33, 118546.05, 172795.67],
 [0.0, 1.0, 0.0, 27892.92, 84710.77, 164470.71],
 [1.0, 0.0, 0.0, 23640.93, 96189.63, 148001.11],
 [0.0, 0.0, 1.0, 15505.73, 127382.3, 35534.17],
 [1.0, 0.0, 0.0, 22177.74, 154806.14, 28334.72],
 [0.0, 0.0, 1.0, 1000.23, 124153.04, 1903.93],
 [0.0, 1.0, 0.0, 1315.46, 115816.21, 297114.46],
 [1.0, 0.0, 0.0, 0.0, 135426.92, 0.0],
 [0.0, 0.0, 1.0, 542.05, 51743.15, 0.0],
 [1.0, 0.0, 0.0, 0.0, 116983.8, 45173.06]], dtype=object)
```

```
In [10]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.2, random_state=0)
```

```
In [11]: from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(xtrain, ytrain)
ypred = linreg.predict(xtest)
```

```
In [12]: linreg.coef_
```

```
Out[12]: array([ 8.66383692e+01, -8.72645791e+02,  7.86007422e+02,  7.73467193e-01,
 3.28845975e-02,  3.66100259e-02])
```

```
In [13]: linreg.intercept_
```

```
Out[13]: 42467.52924852117
```

```
In [14]: dic = {"Acutal Y":ytest, "Predicted Y":ypred}
df =pd.DataFrame(dic)
df
```

```
Out[14]:
```

| | Acutal Y | Predicted Y |
|---|-----------|---------------|
| 0 | 103282.38 | 103015.201598 |
| 1 | 144259.40 | 132582.277608 |
| 2 | 146121.95 | 132447.738452 |
| 3 | 77798.83 | 71976.098513 |
| 4 | 191050.39 | 178537.482211 |
| 5 | 105008.31 | 116161.242302 |
| 6 | 81229.06 | 67851.692097 |
| 7 | 97483.56 | 98791.733747 |
| 8 | 110352.25 | 113969.435330 |
| 9 | 166187.94 | 167921.065696 |

```
In [15]: from sklearn.metrics import r2_score
```

```
In [16]: print(r2_score(ytest, ypred))
```

```
0.9347068473282567
```

```
In [17]: from sklearn.metrics import mean_absolute_error, mean_squared_error
print(f"MAE -: {mean_absolute_error(ytest, ypred)}")
print(f"MSE -: {mean_squared_error(ytest, ypred)}")
print(f"RMSE -: {np.sqrt(mean_squared_error(ytest, ypred))}")
print(f"R-Squared -: {r2_score(ytest, ypred)}")
```

```
MAE -: 7514.293659643097
MSE -: 83502864.03255911
RMSE -: 9137.990152793946
R-Squared -: 0.9347068473282567
```

```
In [ ]:
```

